

# Lab 2: Cats vs Dogs

In this lab, you will train a convolutional neural network to classify an image into one of two classes: "cat" or "dog". The code for the neural networks you train will be written for you, and you are not (yet!) expected to understand all provided code. However, by the end of the lab, you should be able to:

1. Understand at a high level the training loop for a machine learning model.
2. Understand the distinction between training, validation, and test data.
3. The concepts of overfitting and underfitting.
4. Investigate how different hyperparameters, such as learning rate and batch size, affect the success of training.
5. Compare an ANN (aka Multi-Layer Perceptron) with a CNN.

## What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information.

**Do not submit any other files produced by your code.**

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

With Colab, you can export a PDF file using the menu option **File -> Print** and save as PDF file. **Adjust the scaling to ensure that the text is not cutoff at the margins.**

## Colab Link

Include a link to your colab file here

Colab Link: [https://colab.research.google.com/drive/10ZeNA14w4lhVM40Yv3sXqwFcV\\_7oLS\\_1?usp=sharing](https://colab.research.google.com/drive/10ZeNA14w4lhVM40Yv3sXqwFcV_7oLS_1?usp=sharing)

```
import numpy as np
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torch.utils.data.sampler import SubsetRandomSampler
import torchvision.transforms as transforms
```

## Part 0. Helper Functions

We will be making use of the following helper functions. You will be asked to look at and possibly modify some of these, but you are not expected to understand all of them.

You should look at the function names and read the docstrings. If you are curious, come back and explore the code *after* making some progress on the lab.

```
#####  
#####  
# Data Loading  
  
def get_relevant_indices(dataset, classes, target_classes):  
    """ Return the indices for datapoints in the dataset that belongs  
    to the  
        desired target classes, a subset of all possible classes.  
  
    Args:  
        dataset: Dataset object  
        classes: A list of strings denoting the name of each class  
        target_classes: A list of strings denoting the name of desired  
classes  
                        Should be a subset of the 'classes'  
  
    Returns:  
        indices: list of indices that have labels corresponding to one  
of the  
                target classes  
    """  
    indices = []  
    for i in range(len(dataset)):  
        # Check if the label is in the target classes  
        label_index = dataset[i][1] # ex: 3  
        label_class = classes[label_index] # ex: 'cat'  
        if label_class in target_classes:  
            indices.append(i)  
    return indices  
  
def get_data_loader(target_classes, batch_size):  
    """ Loads images of cats and dogs, splits the data into training,  
validation  
        and testing datasets. Returns data loaders for the three  
preprocessed datasets.  
  
    Args:  
        target_classes: A list of strings denoting the name of the  
desired  
                        classes. Should be a subset of the argument  
'classes'  
        batch_size: A int representing the number of samples per batch
```

```

    Returns:
        train_loader: iterable training dataset organized according to
        batch size
        val_loader: iterable validation dataset organized according to
        batch size
        test_loader: iterable testing dataset organized according to
        batch size
        classes: A list of strings denoting the name of each class
    """

    classes = ('plane', 'car', 'bird', 'cat',
               'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

#####
##
# The output of torchvision datasets are PILImage images of range
[0, 1].
# We transform them to Tensors of normalized range [-1, 1].
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
# Load CIFAR10 training data
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True,

transform=transform)
# Get the list of indices to sample from
relevant_indices = get_relevant_indices(trainset, classes,
target_classes)

# Split into train and validation
np.random.seed(1000) # Fixed numpy random seed for reproducible
shuffling
np.random.shuffle(relevant_indices)
split = int(len(relevant_indices) * 0.8) #split at 80%

# split into training and validation indices
relevant_train_indices, relevant_val_indices =
relevant_indices[:split], relevant_indices[split:]
train_sampler = SubsetRandomSampler(relevant_train_indices)
train_loader = torch.utils.data.DataLoader(trainset,
batch_size=batch_size,
                                         num_workers=1,

sampler=train_sampler)
val_sampler = SubsetRandomSampler(relevant_val_indices)
val_loader = torch.utils.data.DataLoader(trainset,
batch_size=batch_size,
                                         num_workers=1,

sampler=val_sampler)
# Load CIFAR10 testing data
testset = torchvision.datasets.CIFAR10(root='./data', train=False,

```

```

                                                    download=True,
transform=transform)
    # Get the list of indices to sample from
    relevant_test_indices = get_relevant_indices(testset, classes,
target_classes)
    test_sampler = SubsetRandomSampler(relevant_test_indices)
    test_loader = torch.utils.data.DataLoader(testset,
batch_size=batch_size,
                                                    num_workers=1,
sampler=test_sampler)
    return train_loader, val_loader, test_loader, classes

#####
#####
# Training
def get_model_name(name, batch_size, learning_rate, epoch):
    """ Generate a name for the model consisting of all the
hyperparameter values

    Args:
        config: Configuration object containing the hyperparameters
    Returns:
        path: A string with the hyperparameter name and value
concatenated
    """
    path = "model_{0}_bs{1}_lr{2}_epoch{3}".format(name,
                                                    batch_size,
                                                    learning_rate,
                                                    epoch)

    return path

def normalize_label(labels):
    """
    Given a tensor containing 2 possible values, normalize this to 0/1

    Args:
        labels: a 1D tensor containing two possible scalar values
    Returns:
        A tensor normalize to 0/1 value
    """
    max_val = torch.max(labels)
    min_val = torch.min(labels)
    norm_labels = (labels - min_val)/(max_val - min_val)
    return norm_labels

def evaluate(net, loader, criterion):
    """ Evaluate the network on the validation set.

    Args:
        net: PyTorch neural network object

```

```

        loader: PyTorch data loader for the validation set
        criterion: The loss function
    Returns:
        err: A scalar for the avg classification error over the
validation set
        loss: A scalar for the average loss function over the
validation set
    """
    total_loss = 0.0
    total_err = 0.0
    total_epoch = 0
    for i, data in enumerate(loader, 0):
        inputs, labels = data
        labels = normalize_label(labels) # Convert labels to 0/1
        outputs = net(inputs)
        loss = criterion(outputs, labels.float())
        corr = (outputs > 0.0).squeeze().long() != labels
        total_err += int(corr.sum())
        total_loss += loss.item()
        total_epoch += len(labels)
    err = float(total_err) / total_epoch
    loss = float(total_loss) / (i + 1)
    return err, loss

#####
#####
# Training Curve
def plot_training_curve(path):
    """ Plots the training curve for a model run, given the csv files
    containing the train/validation error/loss.

    Args:
        path: The base path of the csv files produced during training
    """
    import matplotlib.pyplot as plt
    train_err = np.loadtxt("{}_train_err.csv".format(path))
    val_err = np.loadtxt("{}_val_err.csv".format(path))
    train_loss = np.loadtxt("{}_train_loss.csv".format(path))
    val_loss = np.loadtxt("{}_val_loss.csv".format(path))
    plt.title("Train vs Validation Error")
    n = len(train_err) # number of epochs
    plt.plot(range(1,n+1), train_err, label="Train")
    plt.plot(range(1,n+1), val_err, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Error")
    plt.legend(loc='best')
    plt.show()
    plt.title("Train vs Validation Loss")
    plt.plot(range(1,n+1), train_loss, label="Train")

```

```
plt.plot(range(1,n+1), val_loss, label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(loc='best')
plt.show()
```

## Part 1. Visualizing the Data [7 pt]

We will make use of some of the CIFAR-10 data set, which consists of colour images of size 32x32 pixels belonging to 10 categories. You can find out more about the dataset at <https://www.cs.toronto.edu/~kriz/cifar.html>

For this assignment, we will only be using the cat and dog categories. We have included code that automatically downloads the dataset the first time that the main script is run.

```
# This will download the CIFAR-10 dataset to a folder called "data"
# the first time you run this code.
```

```
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=1) # One image per batch
```

```
Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to
./data/cifar-10-python.tar.gz
```

```
100%|██████████| 170498071/170498071 [00:01<00:00, 105791520.74it/s]
```

```
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified
```

### Part (a) -- 1 pt

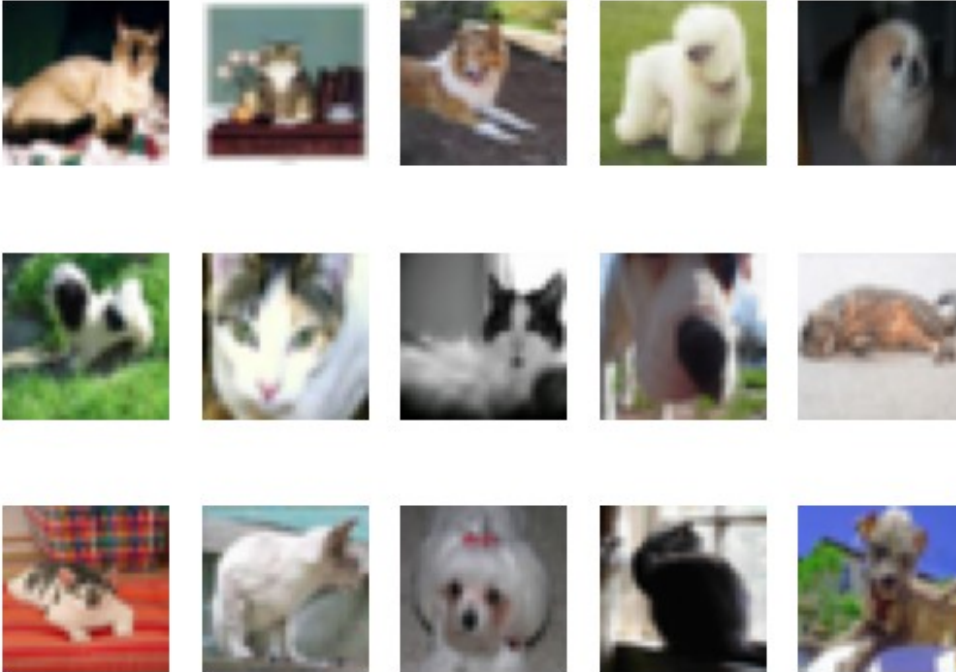
Visualize some of the data by running the code below. Include the visualization in your writeup.

(You don't need to submit anything else.)

```
import matplotlib.pyplot as plt

k = 0
for images, labels in train_loader:
    # since batch_size = 1, there is only 1 image in `images`
    image = images[0]
    # place the colour channel at the end, instead of at the beginning
    img = np.transpose(image, [1,2,0])
    # normalize pixel intensity values to [0, 1]
    img = img / 2 + 0.5
    plt.subplot(3, 5, k+1)
    plt.axis('off')
    plt.imshow(img)
```

```
k += 1
if k > 14:
    break
```



### Part (b) -- 3 pt

How many training examples do we have for the combined `cat` and `dog` classes? What about validation examples? What about test examples?

```
print("Training examples:", len(train_loader.dataset))
print("Validation examples:", len(val_loader.dataset))
print("Test examples:", len(test_loader.dataset))
```

```
Training examples: 50000
Validation examples: 50000
Test examples: 10000
```

### Part (c) -- 3pt

Why do we need a validation set when training our model? What happens if we judge the performance of our models using the training set loss/error instead of the validation set loss/error?

```
# A validation set is needed to test the model on unseen data, helping
to avoid
# overfitting by ensuring the model doesn't just memorize the training
data but
```

```
# learns to generalize. We would also need something to track validation accuracy in the training curve.
```

*# Judging a model only by its performance on the training set can mislead you into thinking it's better than it actually is, as it may not perform well on new, unseen data.*

## Part 2. Training [15 pt]

We define two neural networks, a `LargeNet` and `SmallNet`. We'll be training the networks in this section.

You won't understand fully what these networks are doing until the next few classes, and that's okay. For this assignment, please focus on learning how to train networks, and how hyperparameters affect training.

```
class LargeNet(nn.Module):
    def __init__(self):
        super(LargeNet, self).__init__()
        self.name = "large"
        self.conv1 = nn.Conv2d(3, 5, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(5, 10, 5)
        self.fc1 = nn.Linear(10 * 5 * 5, 32)
        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 10 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        x = x.squeeze(1) # Flatten to [batch_size]
        return x

class SmallNet(nn.Module):
    def __init__(self):
        super(SmallNet, self).__init__()
        self.name = "small"
        self.conv = nn.Conv2d(3, 5, 3)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc = nn.Linear(5 * 7 * 7, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv(x)))
        x = self.pool(x)
```



```

        x = x.view(-1, 5 * 7 * 7)
        x = self.fc(x)
        x = x.squeeze(1) # Flatten to [batch_size]
        return x

small_net = SmallNet()
large_net = LargeNet()

```

## Part (a) -- 2pt

The methods `small_net.parameters()` and `large_net.parameters()` produces an iterator of all the trainable parameters of the network. These parameters are torch tensors containing many scalar values.

We haven't learned how how the parameters in these high-dimensional tensors will be used, but we should be able to count the number of parameters. Measuring the number of parameters in a network is one way of measuring the "size" of a network.

What is the total number of parameters in `small_net` and in `large_net`? (Hint: how many numbers are in each tensor?)

```

for param in small_net.parameters():
    print(param.shape)
# The total number of parameters in small_net is:
# (5*3*3*3)+(5)+(1*245)+(1)=386

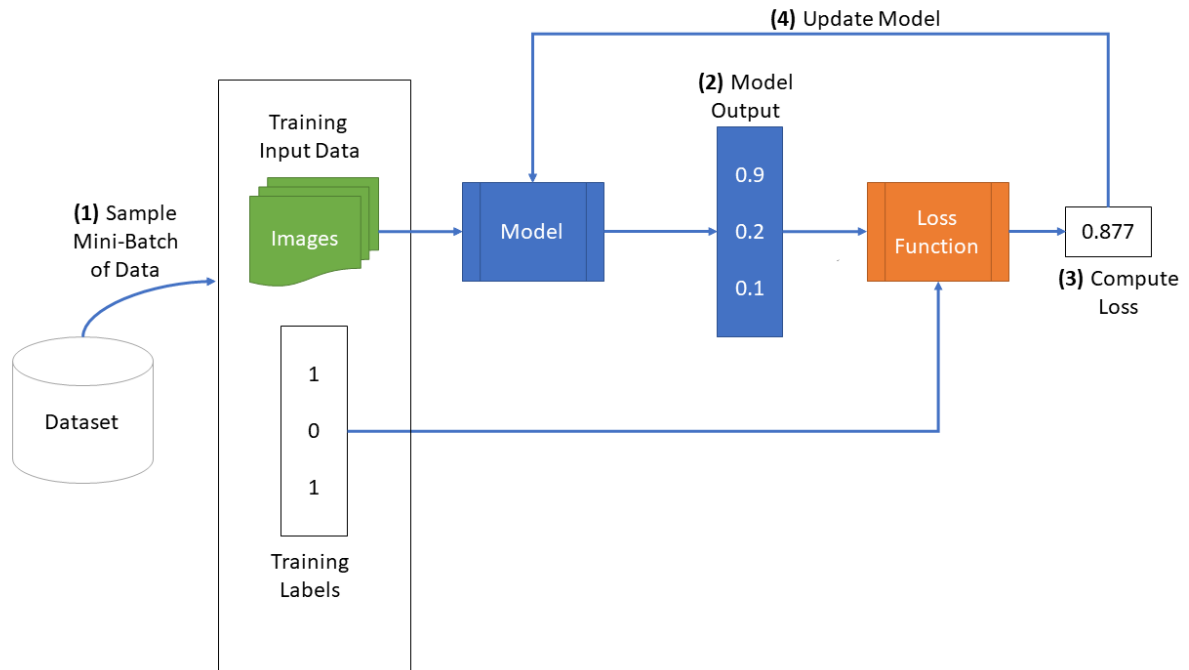
for param in large_net.parameters():
    print(param.shape)
# The total number of parameters in large_net is:
# (5*3*5*5)+(5)+(10*5*5*5)+(10)+(32*250)+(32)+(1*32)+(1)=9705

torch.Size([5, 3, 3, 3])
torch.Size([5])
torch.Size([1, 245])
torch.Size([1])
torch.Size([5, 3, 5, 5])
torch.Size([5])
torch.Size([10, 5, 5, 5])
torch.Size([10])
torch.Size([32, 250])
torch.Size([32])
torch.Size([1, 32])
torch.Size([1])

```

## The function train\_net

The function `train_net` below takes an untrained neural network (like `small_net` and `large_net`) and several other parameters. You should be able to understand how this function works. The figure below shows the high level training loop for a machine learning model:



```
def train_net(net, batch_size=64, learning_rate=0.01, num_epochs=30):
#####
##
# Train a classifier on cats vs dogs
target_classes = ["cat", "dog"]

#####
##
# Fixed PyTorch random seed for reproducible result
torch.manual_seed(1000)

#####
##
# Obtain the PyTorch data loader objects to load batches of the
# datasets
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes, batch_size)

#####
##
# Define the Loss function and optimizer
# The loss function will be Binary Cross Entropy (BCE). In this
# case we
# will use the BCEWithLogitsLoss which takes unnormalized output
# from
# the neural network and scalar label.
# Optimizer will be SGD with Momentum.
```

```

        criterion = nn.BCEWithLogitsLoss()
        optimizer = optim.SGD(net.parameters(), lr=learning_rate,
momentum=0.9)

#####
##
    # Set up some numpy arrays to store the training/test
loss/erruracy
    train_err = np.zeros(num_epochs)
    train_loss = np.zeros(num_epochs)
    val_err = np.zeros(num_epochs)
    val_loss = np.zeros(num_epochs)

#####
##
    # Train the network
    # Loop over the data iterator and sample a new batch of training
data
    # Get the output from the network, and optimize our loss function.
start_time = time.time()
    for epoch in range(num_epochs): # loop over the dataset multiple
times
        total_train_loss = 0.0
        total_train_err = 0.0
        total_epoch = 0
        for i, data in enumerate(train_loader, 0):
            # Get the inputs
            inputs, labels = data
            labels = normalize_label(labels) # Convert labels to 0/1
            # Zero the parameter gradients
            optimizer.zero_grad()
            # Forward pass, backward pass, and optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels.float())
            loss.backward()
            optimizer.step()
            # Calculate the statistics
            corr = (outputs > 0.0).squeeze().long() != labels
            total_train_err += int(corr.sum())
            total_train_loss += loss.item()
            total_epoch += len(labels)
        train_err[epoch] = float(total_train_err) / total_epoch
        train_loss[epoch] = float(total_train_loss) / (i+1)
        val_err[epoch], val_loss[epoch] = evaluate(net, val_loader,
criterion)
        print(("Epoch {}: Train err: {}, Train loss: {} |"+
            "Validation err: {}, Validation loss: {}".format(
                epoch + 1,
                train_err[epoch],

```

```

        train_loss[epoch],
        val_err[epoch],
        val_loss[epoch]))
    # Save the current model (checkpoint) to a file
    model_path = get_model_name(net.name, batch_size,
learning_rate, epoch)
    torch.save(net.state_dict(), model_path)
    print('Finished Training')
    end_time = time.time()
    elapsed_time = end_time - start_time
    print("Total time elapsed: {:.2f} seconds".format(elapsed_time))
    # Write the train/test loss/err into CSV file for plotting later
    epochs = np.arange(1, num_epochs + 1)
    np.savetxt("{}_train_err.csv".format(model_path), train_err)
    np.savetxt("{}_train_loss.csv".format(model_path), train_loss)
    np.savetxt("{}_val_err.csv".format(model_path), val_err)
    np.savetxt("{}_val_loss.csv".format(model_path), val_loss)

```

## Part (b) -- 1pt

The parameters to the function `train_net` are hyperparameters of our neural network. We made these hyperparameters easy to modify so that we can tune them later on.

What are the default values of the parameters `batch_size`, `learning_rate`, and `num_epochs`?

```

# batch_size = 64
# learning_rate = 0.01
# num_epochs = 30

```

## Part (c) -- 3 pt

What files are written to disk when we call `train_net` with `small_net`, and train for 5 epochs? Provide a list of all the files written to disk, and what information the files contain.

```

small_net = SmallNet()
train_net(small_net, num_epochs=5)

# Files written to disk

# model_small_bs64_lr0.01_epoch0: checkpoint after epoch 0
# model_small_bs64_lr0.01_epoch1: checkpoint after epoch 1
# model_small_bs64_lr0.01_epoch2: checkpoint after epoch 2
# model_small_bs64_lr0.01_epoch3: checkpoint after epoch 3
# model_small_bs64_lr0.01_epoch4: checkpoint after epoch 4

# model_small_bs64_lr0.01_epoch4_train_err.csv: file containing
training error
# model_small_bs64_lr0.01_epoch4_train_loss.csv: file containing

```

```
training loss
# model_small_bs64_lr0.01_epoch4_val_err.csv: file containing
validation error
# model_small_bs64_lr0.01_epoch4_val_loss.csv: file containing
validation loss
```

```
Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.43275, Train loss: 0.6750623955726623 |
Validation err: 0.381, Validation loss: 0.6581626161932945
Epoch 2: Train err: 0.371875, Train loss: 0.6496369152069091 |
Validation err: 0.384, Validation loss: 0.6638244986534119
Epoch 3: Train err: 0.35025, Train loss: 0.6355056324005127 |
Validation err: 0.353, Validation loss: 0.6260157078504562
Epoch 4: Train err: 0.339375, Train loss: 0.6214285769462585 |
Validation err: 0.354, Validation loss: 0.624940775334835
Epoch 5: Train err: 0.334125, Train loss: 0.6118814716339112 |
Validation err: 0.3265, Validation loss: 0.6122822929173708
Finished Training
Total time elapsed: 15.39 seconds
```

## Part (d) -- 2pt

Train both `small_net` and `large_net` using the function `train_net` and its default parameters. The function will write many files to disk, including a model checkpoint (saved values of model weights) at the end of each epoch.

If you are using Google Colab, you will need to mount Google Drive so that the files generated by `train_net` gets saved. We will be using these files in part (d). (See the Google Colab tutorial for more information about this.)

Report the total time elapsed when training each network. Which network took longer to train? Why?

```
# Since the function writes files to disk, you will need to mount
# your Google Drive. If you are working on the lab locally, you
# can comment out this code.
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
Mounted at /content/gdrive
```

```
train_net(small_net) # Total time elapsed: 87.89 seconds
train_net(large_net) # Total time elapsed: 97.51 seconds
```

```
# large_net took longer to train because it has more parameters
```

```
Files already downloaded and verified
Files already downloaded and verified
```

Epoch 1: Train err: 0.32075, Train loss: 0.6001654999256134 |  
Validation err: 0.3375, Validation loss: 0.6160770282149315  
Epoch 2: Train err: 0.315375, Train loss: 0.5937726888656616 |  
Validation err: 0.3385, Validation loss: 0.6340117994695902  
Epoch 3: Train err: 0.30825, Train loss: 0.5896465487480164 |  
Validation err: 0.314, Validation loss: 0.606388745829463  
Epoch 4: Train err: 0.311375, Train loss: 0.5856761095523835 |  
Validation err: 0.328, Validation loss: 0.6113690435886383  
Epoch 5: Train err: 0.30225, Train loss: 0.5831401715278626 |  
Validation err: 0.308, Validation loss: 0.6056779511272907  
Epoch 6: Train err: 0.302875, Train loss: 0.5780168986320495 |  
Validation err: 0.3105, Validation loss: 0.6076986379921436  
Epoch 7: Train err: 0.3015, Train loss: 0.5758937051296235 |Validation  
err: 0.323, Validation loss: 0.6030783243477345  
Epoch 8: Train err: 0.301875, Train loss: 0.5725233724117279 |  
Validation err: 0.32, Validation loss: 0.6055176518857479  
Epoch 9: Train err: 0.298, Train loss: 0.5728731179237365 |Validation  
err: 0.328, Validation loss: 0.6091796448454261  
Epoch 10: Train err: 0.294375, Train loss: 0.5680857949256897 |  
Validation err: 0.324, Validation loss: 0.6034962432458997  
Epoch 11: Train err: 0.29575, Train loss: 0.5672789976596833 |  
Validation err: 0.323, Validation loss: 0.6053986735641956  
Epoch 12: Train err: 0.292375, Train loss: 0.5627782278060913 |  
Validation err: 0.316, Validation loss: 0.6011649258434772  
Epoch 13: Train err: 0.294875, Train loss: 0.5666881003379822 |  
Validation err: 0.3215, Validation loss: 0.6052737953141332  
Epoch 14: Train err: 0.288625, Train loss: 0.5624566309452057 |  
Validation err: 0.3305, Validation loss: 0.6153043657541275  
Epoch 15: Train err: 0.29175, Train loss: 0.5615632283687592 |  
Validation err: 0.3285, Validation loss: 0.6142218951135874  
Epoch 16: Train err: 0.292, Train loss: 0.5672653877735138 |Validation  
err: 0.327, Validation loss: 0.6161343650892377  
Epoch 17: Train err: 0.291, Train loss: 0.5630094616413116 |Validation  
err: 0.3215, Validation loss: 0.6008257139474154  
Epoch 18: Train err: 0.290875, Train loss: 0.5610948503017426 |  
Validation err: 0.3275, Validation loss: 0.599173647351563  
Epoch 19: Train err: 0.285375, Train loss: 0.5566081831455231 |  
Validation err: 0.33, Validation loss: 0.6196505427360535  
Epoch 20: Train err: 0.283875, Train loss: 0.5579070656299591 |  
Validation err: 0.3185, Validation loss: 0.6068644756451249  
Epoch 21: Train err: 0.28775, Train loss: 0.5599466969966889 |  
Validation err: 0.317, Validation loss: 0.6008214484900236  
Epoch 22: Train err: 0.289375, Train loss: 0.5601010990142822 |  
Validation err: 0.3245, Validation loss: 0.6129764560610056  
Epoch 23: Train err: 0.284, Train loss: 0.5576876988410949 |Validation  
err: 0.313, Validation loss: 0.6002160776406527  
Epoch 24: Train err: 0.28275, Train loss: 0.5538270087242126 |  
Validation err: 0.3075, Validation loss: 0.6007529767230153  
Epoch 25: Train err: 0.286, Train loss: 0.5545855436325073 |Validation

err: 0.3225, Validation loss: 0.6041560107842088  
Epoch 26: Train err: 0.284625, Train loss: 0.5557803564071655 |  
Validation err: 0.3205, Validation loss: 0.6072382936254144  
Epoch 27: Train err: 0.2805, Train loss: 0.5539706726074218 |  
Validation err: 0.315, Validation loss: 0.6098390519618988  
Epoch 28: Train err: 0.28075, Train loss: 0.5540334072113037 |  
Validation err: 0.3205, Validation loss: 0.5990819809958339  
Epoch 29: Train err: 0.286, Train loss: 0.5562302494049072 |Validation  
err: 0.315, Validation loss: 0.6086240578442812  
Epoch 30: Train err: 0.2775, Train loss: 0.5522372324466706 |  
Validation err: 0.3255, Validation loss: 0.6091442313045263  
Finished Training  
Total time elapsed: 87.89 seconds  
Files already downloaded and verified  
Files already downloaded and verified  
Epoch 1: Train err: 0.455875, Train loss: 0.6898005332946777 |  
Validation err: 0.42, Validation loss: 0.6789227221161127  
Epoch 2: Train err: 0.42075, Train loss: 0.677659417629242 |Validation  
err: 0.433, Validation loss: 0.6792405322194099  
Epoch 3: Train err: 0.401875, Train loss: 0.6654470906257629 |  
Validation err: 0.3815, Validation loss: 0.6485890541225672  
Epoch 4: Train err: 0.368875, Train loss: 0.6469981565475463 |  
Validation err: 0.3705, Validation loss: 0.6400767955929041  
Epoch 5: Train err: 0.356, Train loss: 0.6308781847953796 |Validation  
err: 0.349, Validation loss: 0.6240914296358824  
Epoch 6: Train err: 0.337625, Train loss: 0.6159418187141419 |  
Validation err: 0.331, Validation loss: 0.6148061752319336  
Epoch 7: Train err: 0.32725, Train loss: 0.6027624526023865 |  
Validation err: 0.337, Validation loss: 0.6064695063978434  
Epoch 8: Train err: 0.311625, Train loss: 0.5864601356983185 |  
Validation err: 0.3285, Validation loss: 0.6002039462327957  
Epoch 9: Train err: 0.305, Train loss: 0.580260261297226 |Validation  
err: 0.3245, Validation loss: 0.5950973657891154  
Epoch 10: Train err: 0.295125, Train loss: 0.5654756438732147 |  
Validation err: 0.316, Validation loss: 0.589879828505218  
Epoch 11: Train err: 0.28775, Train loss: 0.5537578790187836 |  
Validation err: 0.3085, Validation loss: 0.5859386436641216  
Epoch 12: Train err: 0.273375, Train loss: 0.539053367137909 |  
Validation err: 0.309, Validation loss: 0.5814298605546355  
Epoch 13: Train err: 0.2735, Train loss: 0.5335126321315765 |  
Validation err: 0.315, Validation loss: 0.5896247252821922  
Epoch 14: Train err: 0.259625, Train loss: 0.5164484293460846 |  
Validation err: 0.3055, Validation loss: 0.5819545490667224  
Epoch 15: Train err: 0.25325, Train loss: 0.5125630114078522 |  
Validation err: 0.2935, Validation loss: 0.5702709592878819  
Epoch 16: Train err: 0.25225, Train loss: 0.5099849009513855 |  
Validation err: 0.2945, Validation loss: 0.5685750562697649  
Epoch 17: Train err: 0.244625, Train loss: 0.49768942093849183 |  
Validation err: 0.2945, Validation loss: 0.5717707984149456

```

Epoch 18: Train err: 0.242625, Train loss: 0.4826437952518463 |
Validation err: 0.308, Validation loss: 0.5783763285726309
Epoch 19: Train err: 0.239125, Train loss: 0.4783159794807434 |
Validation err: 0.294, Validation loss: 0.5751320468261838
Epoch 20: Train err: 0.225625, Train loss: 0.46587873005867003 |
Validation err: 0.2975, Validation loss: 0.5719575732946396
Epoch 21: Train err: 0.2275, Train loss: 0.4578790512084961 |
Validation err: 0.295, Validation loss: 0.5854484271258116
Epoch 22: Train err: 0.212875, Train loss: 0.4426322093009949 |
Validation err: 0.2895, Validation loss: 0.6023599095642567
Epoch 23: Train err: 0.211625, Train loss: 0.439261545419693 |
Validation err: 0.297, Validation loss: 0.5798414368182421
Epoch 24: Train err: 0.201625, Train loss: 0.42281155157089234 |
Validation err: 0.29, Validation loss: 0.5895498618483543
Epoch 25: Train err: 0.192, Train loss: 0.4079967005252838 |Validation
err: 0.2955, Validation loss: 0.6151111563667655
Epoch 26: Train err: 0.183125, Train loss: 0.38918790316581725 |
Validation err: 0.282, Validation loss: 0.6292725810781121
Epoch 27: Train err: 0.178625, Train loss: 0.3870407974720001 |
Validation err: 0.299, Validation loss: 0.6327890455722809
Epoch 28: Train err: 0.175125, Train loss: 0.3769093055725098 |
Validation err: 0.2955, Validation loss: 0.6174397924914956
Epoch 29: Train err: 0.16725, Train loss: 0.3661430332660675 |
Validation err: 0.3115, Validation loss: 0.7061819490045309
Epoch 30: Train err: 0.162, Train loss: 0.35297489547729494 |
Validation err: 0.3, Validation loss: 0.6480715395882726
Finished Training
Total time elapsed: 97.51 seconds

```

## Part (e) - 2pt

Use the function `plot_training_curve` to display the trajectory of the training/validation error and the training/validation loss. You will need to use the function `get_model_name` to generate the argument to the `plot_training_curve` function.

Do this for both the small network and the large network. Include both plots in your writeup.

```

#model_path = get_model_name("small", batch_size=??, learning_rate=??,
epoch=29)

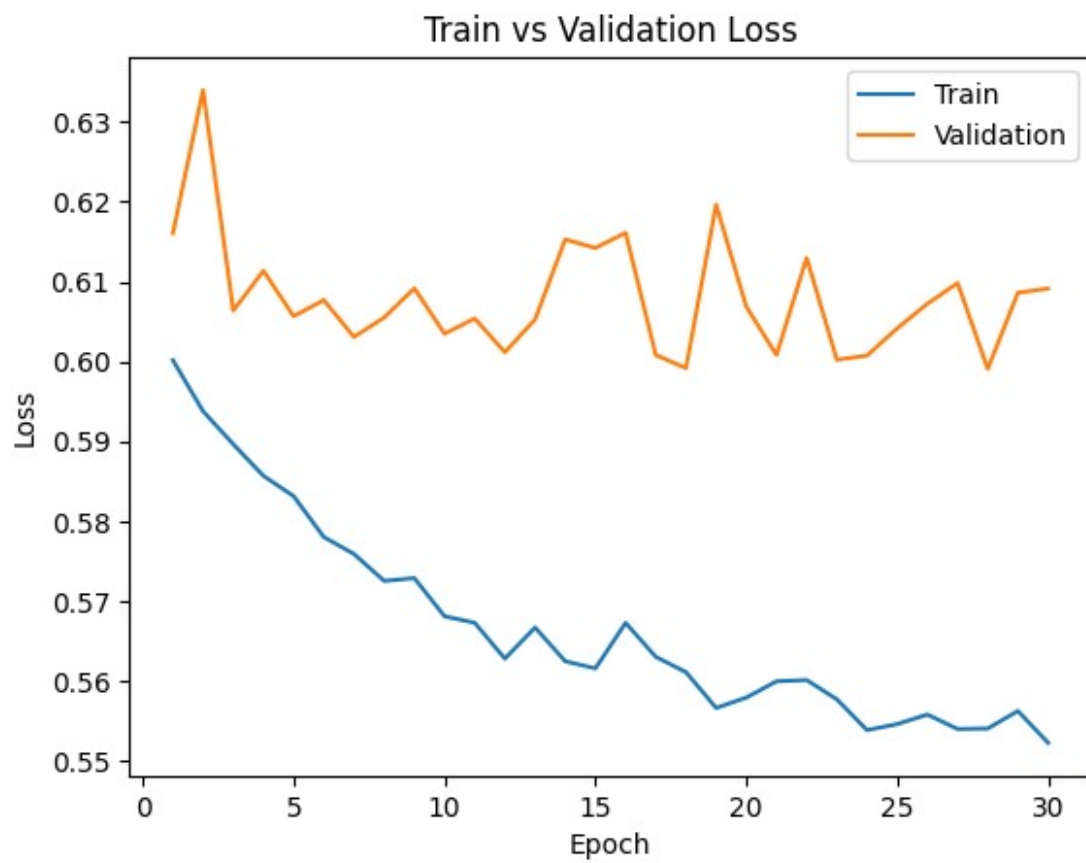
print("Training curve: small_net")
model_path_small = get_model_name("small", batch_size=64,
learning_rate=0.01, epoch=29)
plot_training_curve(model_path_small)
print("Training curve: large_net")
model_path_large = get_model_name("large", batch_size=64,
learning_rate=0.01, epoch=29)
plot_training_curve(model_path_large)

```



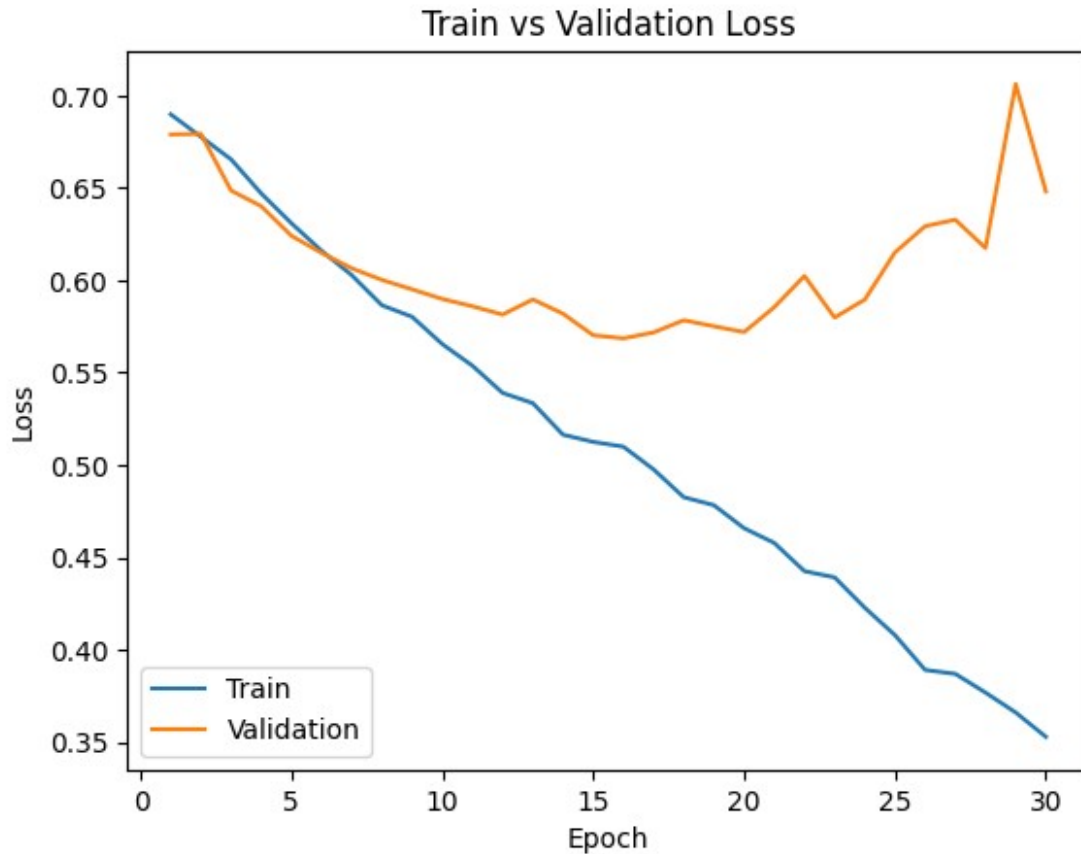
Training curve: small\_net





Training curve: large\_net





## Part (f) - 5pt

Describe what you notice about the training curve. How do the curves differ for `small_net` and `large_net`? Identify any occurrences of underfitting and overfitting.

```
# For small_net, the model shows underfitting near the beginning of
# the curves
# for both validation error and loss graphs. The train and validation
# curves are
# also quite far apart.

# For large_net, the model shows underfitting near the beginning and
# it
# eventually leads to overfitting near the end. The training error and
# loss
# continues decreasing as epoch increases, but validation error and
# loss slowly
# stops decreasing around the half-way point.
```

## Part 3. Optimization Parameters [12 pt]

For this section, we will work with `large_net` only.

## Part (a) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.001`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *lowering* the learning rate.

```
# Note: When we re-construct the model, we start the training
# with *random weights*. If we omit this code, the values of
# the weights will still be the previously trained values.
large_net = LargeNet()
train_net(large_net, 64, 0.001, 30)
model_path_large = get_model_name("large", batch_size=64,
learning_rate=0.001, epoch=29)
plot_training_curve(model_path_large)

# Total time elapsed: 94.83 seconds
# The model takes approximately the same to train.
# Lowering the learning rate does not significantly change the
# training time but
# results in higher error and loss on the training data. For
# validation data,
# while the error is higher, the lower learning rate helps avoid
# overfitting,
# indicating that 0.001 may be too slow for efficient training in this
# case.
```

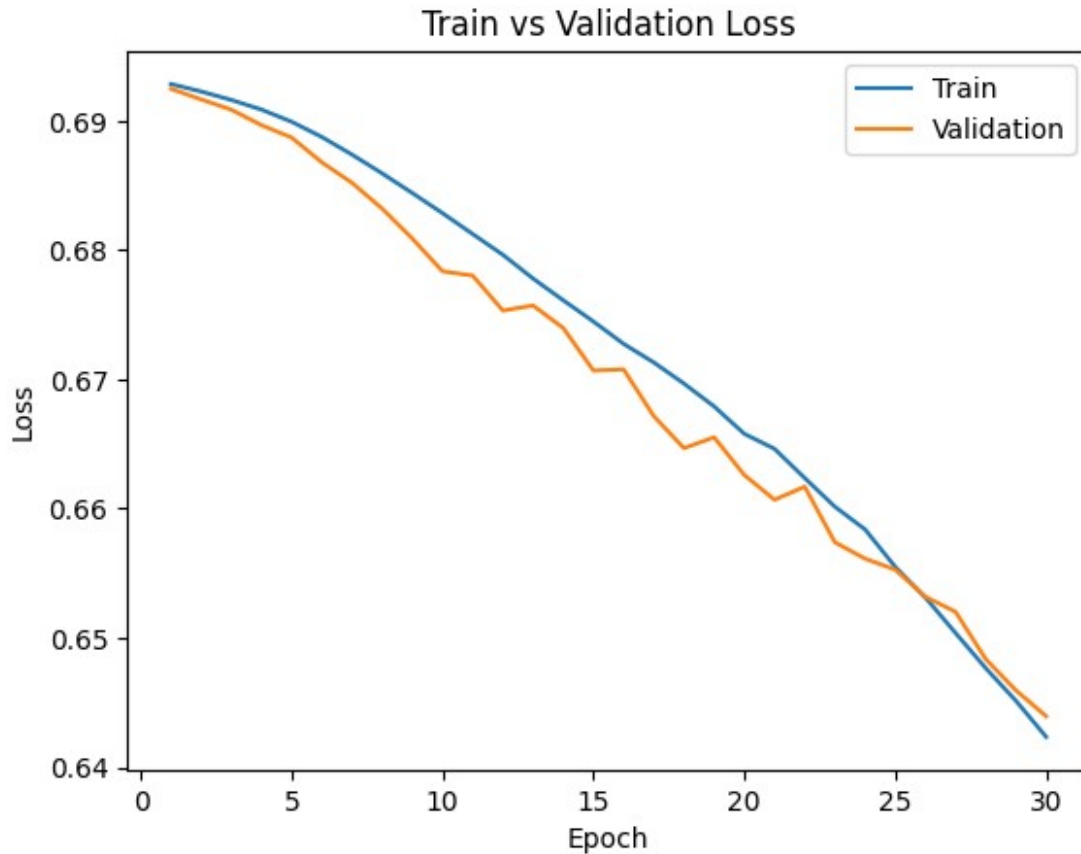
Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.47625, Train loss: 0.6928360028266907 |
Validation err: 0.467, Validation loss: 0.6924686599522829
Epoch 2: Train err: 0.448625, Train loss: 0.6922589740753173 |
Validation err: 0.4305, Validation loss: 0.6916493307799101
Epoch 3: Train err: 0.43575, Train loss: 0.6916067419052124 |
Validation err: 0.4285, Validation loss: 0.6908544395118952
Epoch 4: Train err: 0.430125, Train loss: 0.6908613877296448 |
Validation err: 0.424, Validation loss: 0.6896596923470497
Epoch 5: Train err: 0.434125, Train loss: 0.6899198365211486 |
Validation err: 0.4195, Validation loss: 0.6886942777782679
Epoch 6: Train err: 0.435875, Train loss: 0.6887419753074646 |
Validation err: 0.4195, Validation loss: 0.6867837514728308
Epoch 7: Train err: 0.436625, Train loss: 0.6873781814575195 |
Validation err: 0.4185, Validation loss: 0.6851996649056673
Epoch 8: Train err: 0.43725, Train loss: 0.6859267921447754 |
Validation err: 0.4115, Validation loss: 0.6831991747021675
Epoch 9: Train err: 0.424375, Train loss: 0.6844045443534851 |
Validation err: 0.411, Validation loss: 0.6808868050575256
Epoch 10: Train err: 0.424375, Train loss: 0.682848952293396 |
Validation err: 0.408, Validation loss: 0.6783499345183372
Epoch 11: Train err: 0.425375, Train loss: 0.6812354407310486 |
Validation err: 0.4125, Validation loss: 0.6780234910547733
```

Epoch 12: Train err: 0.419875, Train loss: 0.6796327586174011 |  
Validation err: 0.4125, Validation loss: 0.6753160189837217  
Epoch 13: Train err: 0.414625, Train loss: 0.6777928824424744 |  
Validation err: 0.415, Validation loss: 0.6757084671407938  
Epoch 14: Train err: 0.412, Train loss: 0.6761129403114319 |Validation  
err: 0.412, Validation loss: 0.6739710867404938  
Epoch 15: Train err: 0.409, Train loss: 0.6744727687835693 |Validation  
err: 0.415, Validation loss: 0.6706812102347612  
Epoch 16: Train err: 0.4065, Train loss: 0.67274143409729 |Validation  
err: 0.4105, Validation loss: 0.670768965035677  
Epoch 17: Train err: 0.40125, Train loss: 0.6713059720993042 |  
Validation err: 0.404, Validation loss: 0.66715326346457  
Epoch 18: Train err: 0.399375, Train loss: 0.6696724286079406 |  
Validation err: 0.4055, Validation loss: 0.6646785754710436  
Epoch 19: Train err: 0.40075, Train loss: 0.6679023985862732 |  
Validation err: 0.3955, Validation loss: 0.6655161324888468  
Epoch 20: Train err: 0.3925, Train loss: 0.6657856888771058 |  
Validation err: 0.405, Validation loss: 0.6625996753573418  
Epoch 21: Train err: 0.38975, Train loss: 0.6646266269683838 |  
Validation err: 0.395, Validation loss: 0.6606872715055943  
Epoch 22: Train err: 0.38875, Train loss: 0.6623701963424683 |  
Validation err: 0.3935, Validation loss: 0.6617010589689016  
Epoch 23: Train err: 0.384125, Train loss: 0.6601490645408631 |  
Validation err: 0.3975, Validation loss: 0.6573988571763039  
Epoch 24: Train err: 0.3825, Train loss: 0.6583953781127929 |  
Validation err: 0.386, Validation loss: 0.6561295725405216  
Epoch 25: Train err: 0.3785, Train loss: 0.6554959454536438 |  
Validation err: 0.3875, Validation loss: 0.6552845854312181  
Epoch 26: Train err: 0.376625, Train loss: 0.6531237239837646 |  
Validation err: 0.387, Validation loss: 0.6531836222857237  
Epoch 27: Train err: 0.37525, Train loss: 0.6503791484832764 |  
Validation err: 0.3875, Validation loss: 0.652014534920454  
Epoch 28: Train err: 0.371375, Train loss: 0.6476589822769165 |  
Validation err: 0.3875, Validation loss: 0.6483639199286699  
Epoch 29: Train err: 0.36825, Train loss: 0.6451378240585327 |  
Validation err: 0.382, Validation loss: 0.6459472719579935  
Epoch 30: Train err: 0.362625, Train loss: 0.6423516173362732 |  
Validation err: 0.379, Validation loss: 0.6439454797655344  
Finished Training  
Total time elapsed: 94.83 seconds





## Part (b) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.1`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the learning rate.

```
large_net = LargeNet()
train_net(large_net, 64, 0.1, 30)
model_path_large = get_model_name("large", batch_size=64,
learning_rate=0.1, epoch=29)
plot_training_curve(model_path_large)
```

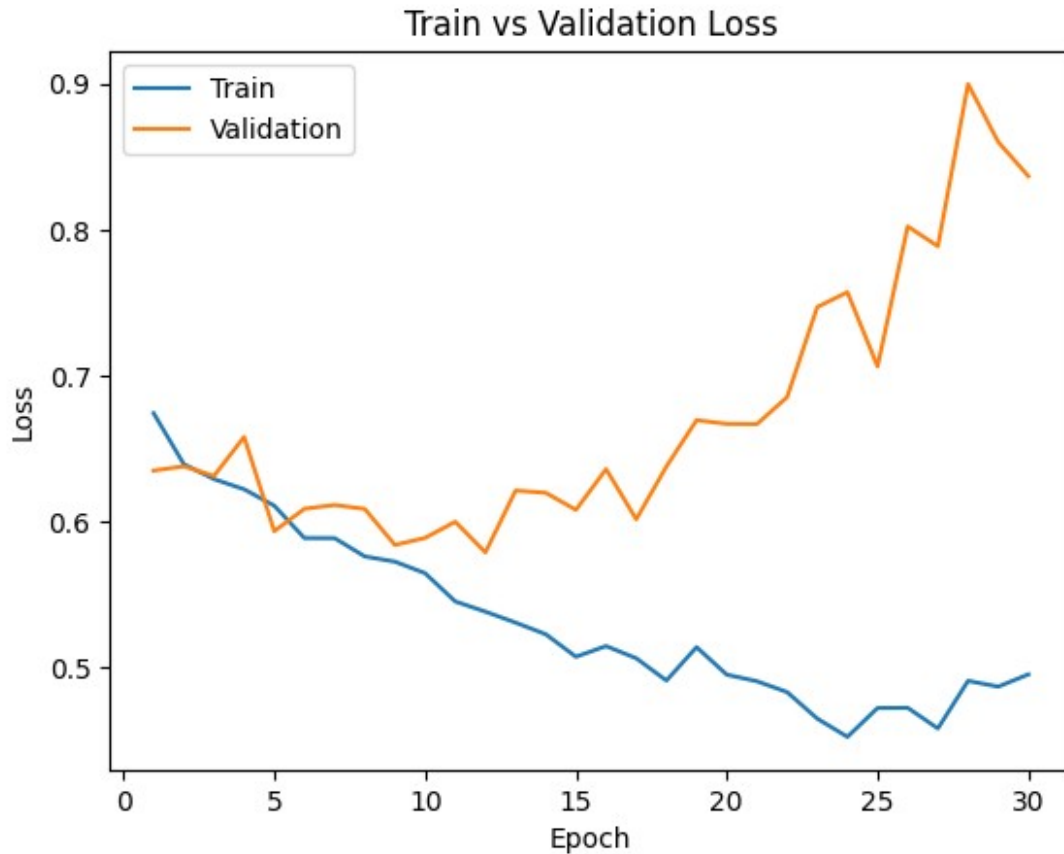
*# Total time elapsed: 94.88 seconds*  
*# The model takes approximately the same to train.*  
*# Increasing the learning rate does not significantly change the training time*  
*# but results in slightly higher error and loss on the training data.*  
*For*  
*# validation data, it increases around the half-way point indicated overfitting.*  
*# This indicates that 0.1 may be too fast.*



Files already downloaded and verified  
Files already downloaded and verified  
Epoch 1: Train err: 0.4295, Train loss: 0.6743808498382569 | Validation  
err: 0.3565, Validation loss: 0.6350402869284153  
Epoch 2: Train err: 0.364375, Train loss: 0.6394594783782959 |  
Validation err: 0.364, Validation loss: 0.6380155384540558  
Epoch 3: Train err: 0.357125, Train loss: 0.6292509632110596 |  
Validation err: 0.357, Validation loss: 0.6313505992293358  
Epoch 4: Train err: 0.351375, Train loss: 0.6222558770179748 |  
Validation err: 0.3615, Validation loss: 0.6581139843910933  
Epoch 5: Train err: 0.3425, Train loss: 0.6109966540336609 | Validation  
err: 0.327, Validation loss: 0.593532383441925  
Epoch 6: Train err: 0.31975, Train loss: 0.5888414912223816 |  
Validation err: 0.3045, Validation loss: 0.6088876193389297  
Epoch 7: Train err: 0.31825, Train loss: 0.5887727000713349 |  
Validation err: 0.3255, Validation loss: 0.6114899981766939  
Epoch 8: Train err: 0.310125, Train loss: 0.5762689855098725 |  
Validation err: 0.358, Validation loss: 0.6087798178195953  
Epoch 9: Train err: 0.304625, Train loss: 0.5726798589229584 |  
Validation err: 0.32, Validation loss: 0.5841292049735785  
Epoch 10: Train err: 0.297125, Train loss: 0.5647666845321655 |  
Validation err: 0.3185, Validation loss: 0.5889163985848427  
Epoch 11: Train err: 0.282, Train loss: 0.5454324214458466 | Validation  
err: 0.3085, Validation loss: 0.600019222125411  
Epoch 12: Train err: 0.278625, Train loss: 0.5384270005226135 |  
Validation err: 0.2995, Validation loss: 0.5790340006351471  
Epoch 13: Train err: 0.274375, Train loss: 0.5309341120719909 |  
Validation err: 0.308, Validation loss: 0.6213697791099548  
Epoch 14: Train err: 0.26975, Train loss: 0.5229758756160736 |  
Validation err: 0.308, Validation loss: 0.6198801156133413  
Epoch 15: Train err: 0.257625, Train loss: 0.5075466320514679 |  
Validation err: 0.312, Validation loss: 0.6081355484202504  
Epoch 16: Train err: 0.25625, Train loss: 0.5148655683994293 |  
Validation err: 0.327, Validation loss: 0.6360824098810554  
Epoch 17: Train err: 0.25475, Train loss: 0.5065258400440216 |  
Validation err: 0.3175, Validation loss: 0.601697607897222  
Epoch 18: Train err: 0.24575, Train loss: 0.49124235701560975 |  
Validation err: 0.3205, Validation loss: 0.6378643093630672  
Epoch 19: Train err: 0.259375, Train loss: 0.5141446299552918 |  
Validation err: 0.3385, Validation loss: 0.6695924252271652  
Epoch 20: Train err: 0.24875, Train loss: 0.4952861943244934 |  
Validation err: 0.3445, Validation loss: 0.6670713275671005  
Epoch 21: Train err: 0.243125, Train loss: 0.4907743418216705 |  
Validation err: 0.334, Validation loss: 0.6669150963425636  
Epoch 22: Train err: 0.23875, Train loss: 0.48334857559204103 |  
Validation err: 0.311, Validation loss: 0.6853150613605976  
Epoch 23: Train err: 0.223625, Train loss: 0.4652157464027405 |  
Validation err: 0.3145, Validation loss: 0.7470605578273535  
Epoch 24: Train err: 0.222875, Train loss: 0.4526974799633026 |  
Validation err: 0.33, Validation loss: 0.7572380006313324

Epoch 25: Train err: 0.2345, Train loss: 0.4724377233982086 |  
Validation err: 0.334, Validation loss: 0.7064428050071001  
Epoch 26: Train err: 0.234625, Train loss: 0.4725676612854004 |  
Validation err: 0.368, Validation loss: 0.8023080676794052  
Epoch 27: Train err: 0.22175, Train loss: 0.45844222140312196 |  
Validation err: 0.357, Validation loss: 0.7886781524866819  
Epoch 28: Train err: 0.237375, Train loss: 0.49102354764938355 |  
Validation err: 0.3375, Validation loss: 0.8998023774474859  
Epoch 29: Train err: 0.237125, Train loss: 0.4870765597820282 |  
Validation err: 0.3385, Validation loss: 0.8601288255304098  
Epoch 30: Train err: 0.239375, Train loss: 0.4954684257507324 |  
Validation err: 0.358, Validation loss: 0.8366707749664783  
Finished Training  
Total time elapsed: 94.88 seconds





### Part (c) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=512`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the batch size.

```
large_net = LargeNet()
train_net(large_net, 512, 0.01, 30)
model_path_large = get_model_name("large", batch_size=512,
learning_rate=0.01, epoch=29)
plot_training_curve(model_path_large)
```

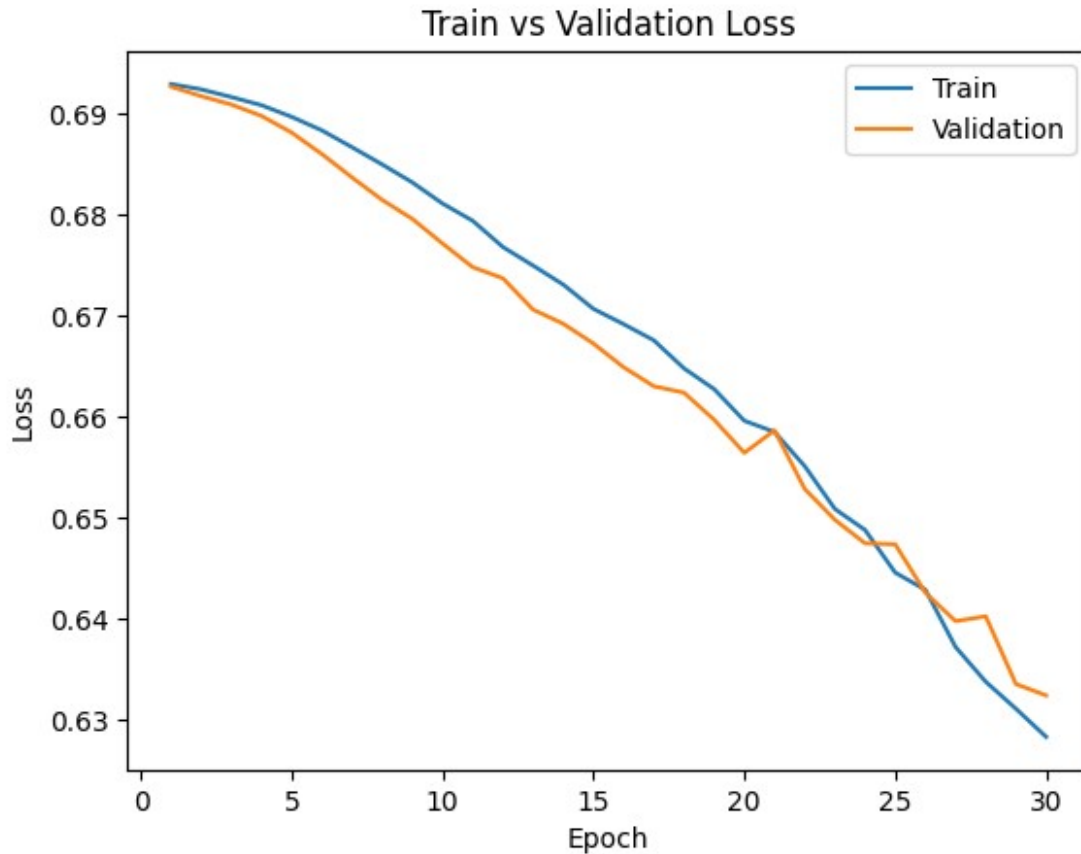
*# Total time elapsed: 82.25 seconds*  
*# The model takes shorter time to train.*  
*# The increase in batch size means the number of iterations within each epoch is*  
*# smaller, reducing the training time. Increasing the batch size eliminated the*  
*# signs of overfitting.*

Files already downloaded and verified  
Files already downloaded and verified  
Epoch 1: Train err: 0.48175, Train loss: 0.6929379515349865 |

Validation err: 0.478, Validation loss: 0.6926823854446411  
Epoch 2: Train err: 0.457625, Train loss: 0.6924104019999504 |  
Validation err: 0.434, Validation loss: 0.6917425245046616  
Epoch 3: Train err: 0.437, Train loss: 0.6916500590741634 |Validation  
err: 0.4265, Validation loss: 0.6909129917621613  
Epoch 4: Train err: 0.433625, Train loss: 0.6908449940383434 |  
Validation err: 0.424, Validation loss: 0.6897870153188705  
Epoch 5: Train err: 0.434, Train loss: 0.6896935552358627 |Validation  
err: 0.424, Validation loss: 0.6881355047225952  
Epoch 6: Train err: 0.438, Train loss: 0.6883531995117664 |Validation  
err: 0.4285, Validation loss: 0.686011865735054  
Epoch 7: Train err: 0.439375, Train loss: 0.6866871826350689 |  
Validation err: 0.426, Validation loss: 0.6836968660354614  
Epoch 8: Train err: 0.43525, Train loss: 0.6849770732223988 |  
Validation err: 0.411, Validation loss: 0.6814672648906708  
Epoch 9: Train err: 0.42375, Train loss: 0.6832008771598339 |  
Validation err: 0.414, Validation loss: 0.6795914322137833  
Epoch 10: Train err: 0.421, Train loss: 0.6811087839305401 |Validation  
err: 0.416, Validation loss: 0.677154079079628  
Epoch 11: Train err: 0.42075, Train loss: 0.6794031001627445 |  
Validation err: 0.4095, Validation loss: 0.6748126447200775  
Epoch 12: Train err: 0.41475, Train loss: 0.6768064312636852 |  
Validation err: 0.412, Validation loss: 0.6737105250358582  
Epoch 13: Train err: 0.410375, Train loss: 0.6749707795679569 |  
Validation err: 0.412, Validation loss: 0.6706106513738632  
Epoch 14: Train err: 0.407125, Train loss: 0.6730907596647739 |  
Validation err: 0.4125, Validation loss: 0.6692123562097549  
Epoch 15: Train err: 0.400375, Train loss: 0.6706820242106915 |  
Validation err: 0.4105, Validation loss: 0.6672529578208923  
Epoch 16: Train err: 0.397625, Train loss: 0.6691752374172211 |  
Validation err: 0.405, Validation loss: 0.6649040132761002  
Epoch 17: Train err: 0.39375, Train loss: 0.6675728745758533 |  
Validation err: 0.4015, Validation loss: 0.6630261093378067  
Epoch 18: Train err: 0.392625, Train loss: 0.6647983901202679 |  
Validation err: 0.394, Validation loss: 0.6623890697956085  
Epoch 19: Train err: 0.386, Train loss: 0.6627401672303677 |Validation  
err: 0.388, Validation loss: 0.6597096621990204  
Epoch 20: Train err: 0.381625, Train loss: 0.6596098616719246 |  
Validation err: 0.4015, Validation loss: 0.6564301252365112  
Epoch 21: Train err: 0.386125, Train loss: 0.6584997698664665 |  
Validation err: 0.3885, Validation loss: 0.6586556434631348  
Epoch 22: Train err: 0.378625, Train loss: 0.6551188267767429 |  
Validation err: 0.385, Validation loss: 0.6528644561767578  
Epoch 23: Train err: 0.371625, Train loss: 0.6508826948702335 |  
Validation err: 0.3835, Validation loss: 0.649808794260025  
Epoch 24: Train err: 0.37675, Train loss: 0.6487980298697948 |  
Validation err: 0.384, Validation loss: 0.6474767625331879  
Epoch 25: Train err: 0.368375, Train loss: 0.6445756293833256 |  
Validation err: 0.382, Validation loss: 0.6473795771598816

Epoch 26: Train err: 0.372875, Train loss: 0.6428777538239956 |  
Validation err: 0.3755, Validation loss: 0.6425858736038208  
Epoch 27: Train err: 0.35925, Train loss: 0.6372104585170746 |  
Validation err: 0.3785, Validation loss: 0.6397744864225388  
Epoch 28: Train err: 0.35425, Train loss: 0.6337734051048756 |  
Validation err: 0.37, Validation loss: 0.6402499973773956  
Epoch 29: Train err: 0.3535, Train loss: 0.6310990080237389 |  
Validation err: 0.366, Validation loss: 0.6335441172122955  
Epoch 30: Train err: 0.353, Train loss: 0.6283206455409527 | Validation  
err: 0.367, Validation loss: 0.6324219256639481  
Finished Training  
Total time elapsed: 82.25 seconds





### Part (d) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=16`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *decreasing* the batch size.

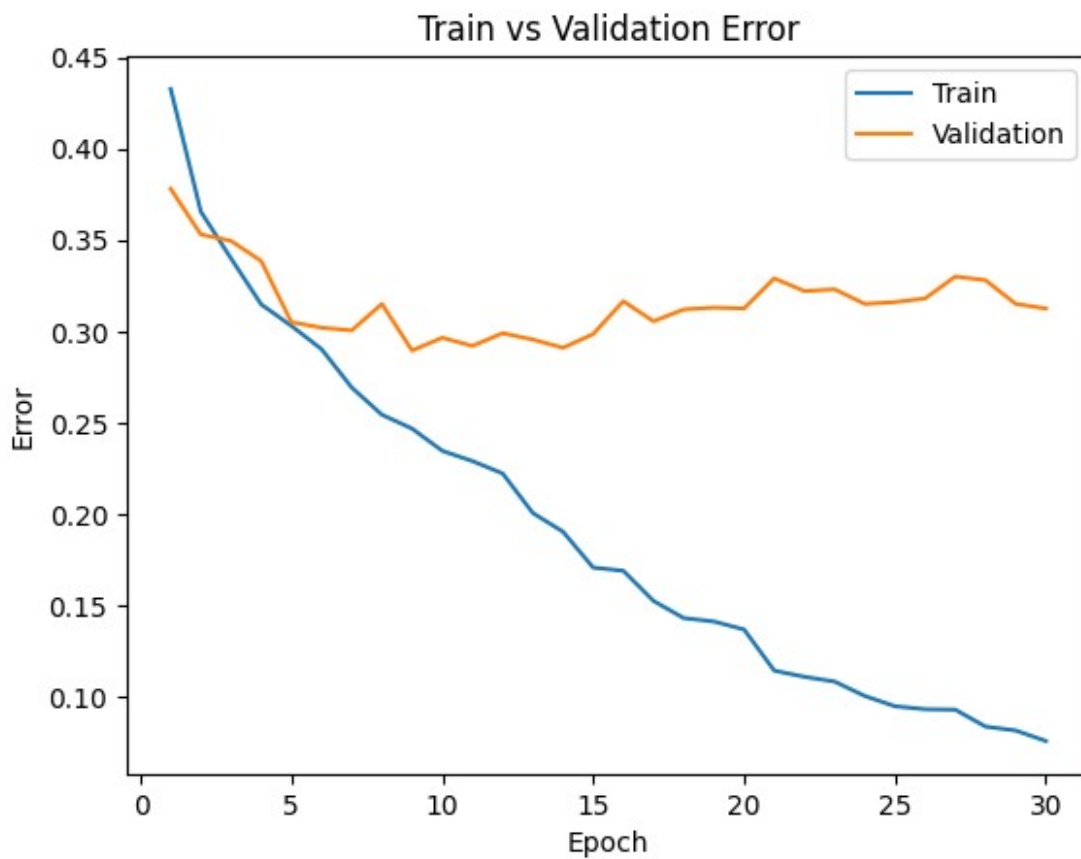
```
large_net = LargeNet()
train_net(large_net, 16, 0.01, 30)
model_path_large = get_model_name("large", batch_size=16,
learning_rate=0.01, epoch=29)
plot_training_curve(model_path_large)
```

*# Total time elapsed: 133.49 seconds*  
*# The model takes longer time to train.*  
*# The decrease in batch size means the number of iterations within*  
*each epoch is*  
*# larger, increasing the training time. The validation error and loss*  
*both show*  
*# overfitting.*

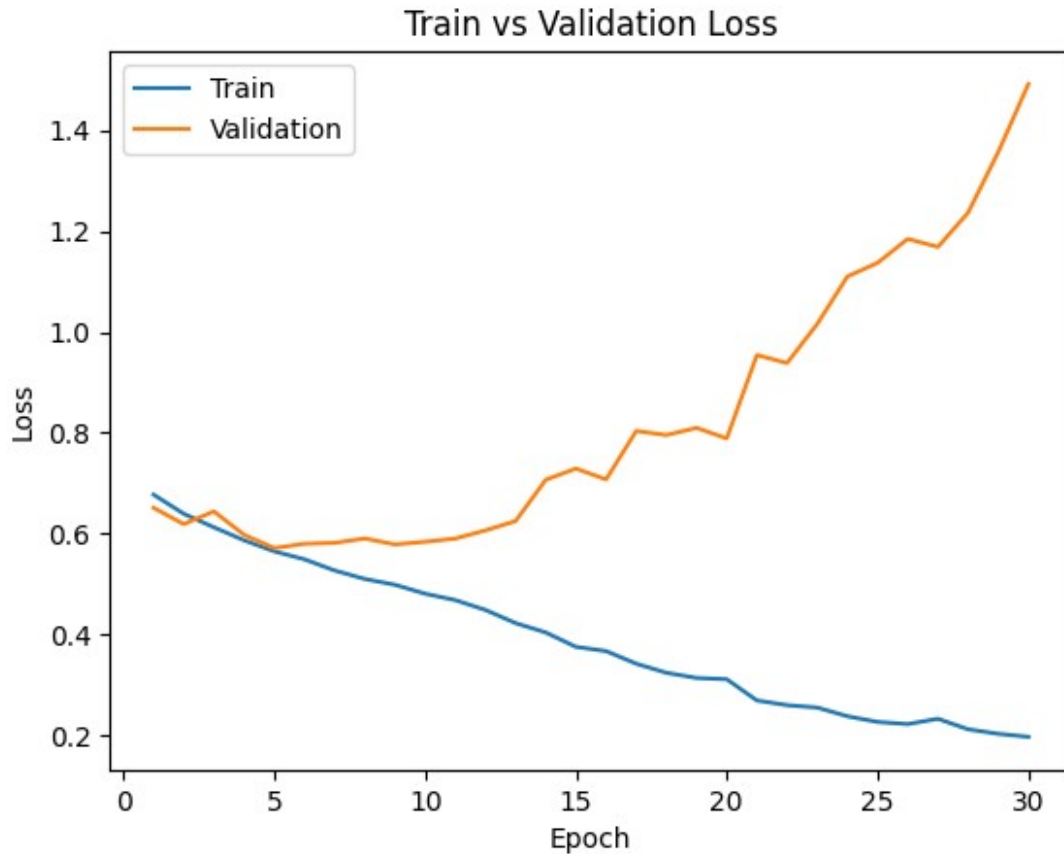
Files already downloaded and verified  
Files already downloaded and verified  
Epoch 1: Train err: 0.432625, Train loss: 0.6775506158471107 |

Validation err: 0.378, Validation loss: 0.6512571625709533  
Epoch 2: Train err: 0.3655, Train loss: 0.6395608481168747 | Validation  
err: 0.353, Validation loss: 0.6191524927616119  
Epoch 3: Train err: 0.339875, Train loss: 0.6127814228534698 |  
Validation err: 0.3495, Validation loss: 0.6442944076061249  
Epoch 4: Train err: 0.31475, Train loss: 0.5870258185863495 |  
Validation err: 0.3385, Validation loss: 0.5978640830516815  
Epoch 5: Train err: 0.303, Train loss: 0.5651368154883385 | Validation  
err: 0.305, Validation loss: 0.5715659594535828  
Epoch 6: Train err: 0.29025, Train loss: 0.5493593747615815 |  
Validation err: 0.302, Validation loss: 0.5801575357913971  
Epoch 7: Train err: 0.26925, Train loss: 0.5272505192756652 |  
Validation err: 0.3005, Validation loss: 0.5818900580406189  
Epoch 8: Train err: 0.254375, Train loss: 0.5104012369215488 |  
Validation err: 0.315, Validation loss: 0.5906455669403076  
Epoch 9: Train err: 0.24675, Train loss: 0.4989313832819462 |  
Validation err: 0.2895, Validation loss: 0.5788275961875915  
Epoch 10: Train err: 0.234625, Train loss: 0.4812071683704853 |  
Validation err: 0.2965, Validation loss: 0.58406099319458  
Epoch 11: Train err: 0.229, Train loss: 0.4685571998953819 | Validation  
err: 0.292, Validation loss: 0.5905215789079666  
Epoch 12: Train err: 0.22225, Train loss: 0.44902220405638216 |  
Validation err: 0.299, Validation loss: 0.6065930938720703  
Epoch 13: Train err: 0.200625, Train loss: 0.4229303933084011 |  
Validation err: 0.2955, Validation loss: 0.6251139466762543  
Epoch 14: Train err: 0.190375, Train loss: 0.4043408792465925 |  
Validation err: 0.291, Validation loss: 0.7069108620882034  
Epoch 15: Train err: 0.17075, Train loss: 0.3758661460578442 |  
Validation err: 0.2985, Validation loss: 0.7291838767528533  
Epoch 16: Train err: 0.169, Train loss: 0.36747281277179716 |  
Validation err: 0.3165, Validation loss: 0.7076600217819213  
Epoch 17: Train err: 0.1525, Train loss: 0.34249910655617716 |  
Validation err: 0.3055, Validation loss: 0.8037151482105255  
Epoch 18: Train err: 0.143125, Train loss: 0.3244629814103246 |  
Validation err: 0.312, Validation loss: 0.7955756268501282  
Epoch 19: Train err: 0.14125, Train loss: 0.31422823867201805 |  
Validation err: 0.313, Validation loss: 0.8098329643011093  
Epoch 20: Train err: 0.136875, Train loss: 0.3117204641513526 |  
Validation err: 0.3125, Validation loss: 0.788938661813736  
Epoch 21: Train err: 0.114375, Train loss: 0.2696312564909458 |  
Validation err: 0.329, Validation loss: 0.9542041051387787  
Epoch 22: Train err: 0.111, Train loss: 0.26029947843030093 |  
Validation err: 0.322, Validation loss: 0.9382650710344315  
Epoch 23: Train err: 0.108375, Train loss: 0.2555502322986722 |  
Validation err: 0.323, Validation loss: 1.0157140600681305  
Epoch 24: Train err: 0.1005, Train loss: 0.23835007748007775 |  
Validation err: 0.315, Validation loss: 1.109861346244812  
Epoch 25: Train err: 0.094875, Train loss: 0.2270894316267222 |  
Validation err: 0.316, Validation loss: 1.1367672670185567

Epoch 26: Train err: 0.09325, Train loss: 0.2228318989351392 |  
Validation err: 0.318, Validation loss: 1.1843619964122771  
Epoch 27: Train err: 0.093, Train loss: 0.23337606593593954 |  
Validation err: 0.33, Validation loss: 1.1684871008992195  
Epoch 28: Train err: 0.08375, Train loss: 0.21251985063403844 |  
Validation err: 0.328, Validation loss: 1.2360798382759095  
Epoch 29: Train err: 0.081625, Train loss: 0.2035546142552048 |  
Validation err: 0.315, Validation loss: 1.356388526082039  
Epoch 30: Train err: 0.075875, Train loss: 0.19755484854709357 |  
Validation err: 0.3125, Validation loss: 1.4912220520973205  
Finished Training  
Total time elapsed: 133.49 seconds







## Part 4. Hyperparameter Search [6 pt]

### Part (a) - 2pt

Based on the plots from above, choose another set of values for the hyperparameters (network, batch\_size, learning\_rate) that you think would help you improve the validation accuracy. Justify your choice.

```
# Hyperparameters were chosen as:  
# network = large_net  
# batch_size = 256  
# learning_rate = 0.01  
# epoch = 30  
# In the parts above, we saw that:  
# large_net works better than small_net  
# 0.01 was the better learning rate between 0.001, 0.01, and 0.1  
# Smaller batch_size resulted in overfitting, therefore having a  
larger  
# batch_size would result in better curves  
# The number of epochs was kept the same
```

## Part (b) - 1pt

Train the model with the hyperparameters you chose in part(a), and include the training curve.

```
large_net = LargeNet()
train_net(large_net, 256, 0.01, 30)
model_path_large = get_model_name("large", batch_size=256,
learning_rate=0.01, epoch=29)
plot_training_curve(model_path_large)
```

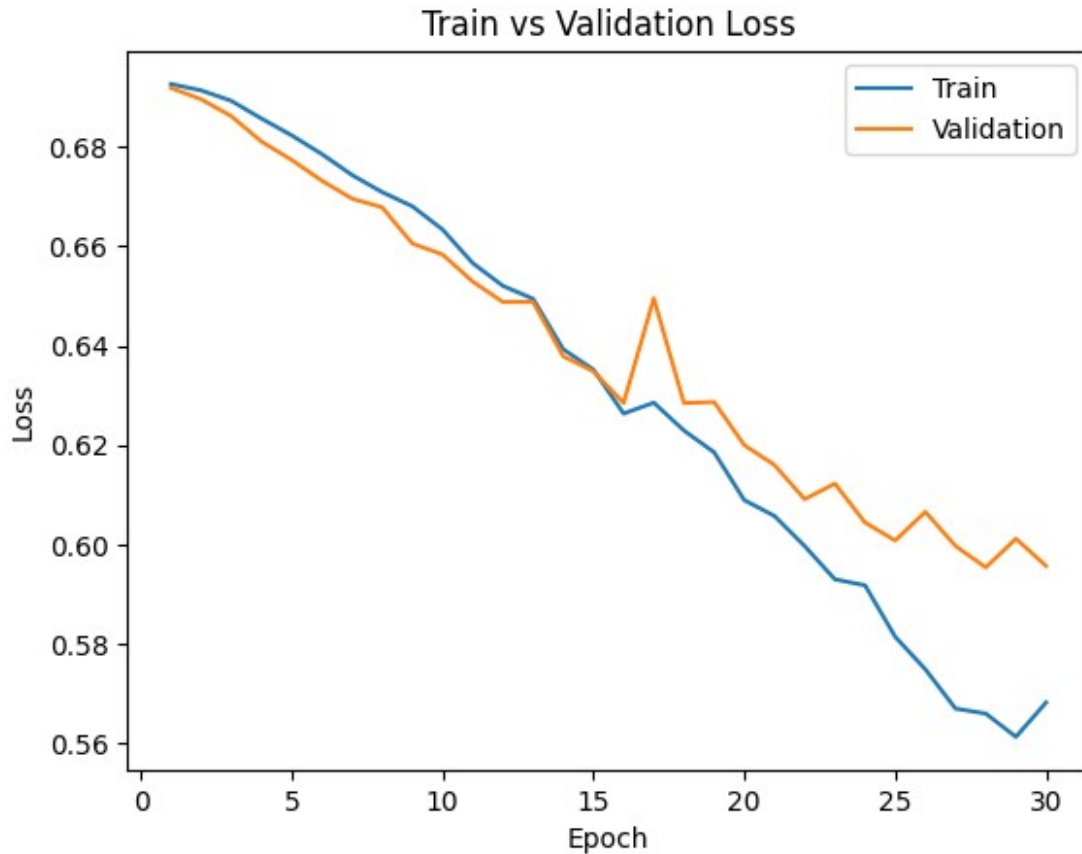
Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.467625, Train loss: 0.6926687732338905 |
Validation err: 0.4355, Validation loss: 0.6918838396668434
Epoch 2: Train err: 0.45125, Train loss: 0.6913921386003494 |
Validation err: 0.446, Validation loss: 0.689614325761795
Epoch 3: Train err: 0.428375, Train loss: 0.6892956364899874 |
Validation err: 0.418, Validation loss: 0.6862076073884964
Epoch 4: Train err: 0.432375, Train loss: 0.6857246831059456 |
Validation err: 0.4185, Validation loss: 0.6810905113816261
Epoch 5: Train err: 0.425875, Train loss: 0.6823424454778433 |
Validation err: 0.414, Validation loss: 0.6773988157510757
Epoch 6: Train err: 0.41775, Train loss: 0.6785747222602367 |
Validation err: 0.415, Validation loss: 0.6732195988297462
Epoch 7: Train err: 0.4085, Train loss: 0.674360191449523 |Validation
err: 0.412, Validation loss: 0.669593334197998
Epoch 8: Train err: 0.401625, Train loss: 0.6709097884595394 |
Validation err: 0.4015, Validation loss: 0.6678433567285538
Epoch 9: Train err: 0.398625, Train loss: 0.6680476255714893 |
Validation err: 0.404, Validation loss: 0.6605530977249146
Epoch 10: Train err: 0.38875, Train loss: 0.6633832901716232 |
Validation err: 0.3985, Validation loss: 0.6583485454320908
Epoch 11: Train err: 0.378875, Train loss: 0.6566182859241962 |
Validation err: 0.3865, Validation loss: 0.6529050841927528
Epoch 12: Train err: 0.375, Train loss: 0.6520354263484478 |Validation
err: 0.385, Validation loss: 0.6487870514392853
Epoch 13: Train err: 0.373375, Train loss: 0.6494291331619024 |
Validation err: 0.382, Validation loss: 0.6488054469227791
Epoch 14: Train err: 0.35575, Train loss: 0.6392552871257067 |
Validation err: 0.369, Validation loss: 0.637849822640419
Epoch 15: Train err: 0.3585, Train loss: 0.6351705007255077 |
Validation err: 0.363, Validation loss: 0.6348292082548141
Epoch 16: Train err: 0.3515, Train loss: 0.6263151150196791 |
Validation err: 0.363, Validation loss: 0.6284611001610756
Epoch 17: Train err: 0.3485, Train loss: 0.6285038739442825 |
Validation err: 0.3875, Validation loss: 0.6495563760399818
Epoch 18: Train err: 0.346, Train loss: 0.6229071728885174 |Validation
err: 0.3555, Validation loss: 0.628441222012043
Epoch 19: Train err: 0.348375, Train loss: 0.6185130663216114 |
Validation err: 0.3585, Validation loss: 0.628648966550827
```

```
Epoch 20: Train err: 0.339375, Train loss: 0.608886506408453 |  
Validation err: 0.342, Validation loss: 0.6199510842561722  
Epoch 21: Train err: 0.331875, Train loss: 0.6057094093412161 |  
Validation err: 0.3365, Validation loss: 0.6159674450755119  
Epoch 22: Train err: 0.328625, Train loss: 0.5996328052133322 |  
Validation err: 0.3285, Validation loss: 0.6090726181864738  
Epoch 23: Train err: 0.32375, Train loss: 0.59294911660254 |Validation  
err: 0.3335, Validation loss: 0.6122196018695831  
Epoch 24: Train err: 0.32425, Train loss: 0.5917343720793724 |  
Validation err: 0.337, Validation loss: 0.6043870896100998  
Epoch 25: Train err: 0.30925, Train loss: 0.5813672374933958 |  
Validation err: 0.3255, Validation loss: 0.6007594466209412  
Epoch 26: Train err: 0.307875, Train loss: 0.5748253837227821 |  
Validation err: 0.3255, Validation loss: 0.6065291538834572  
Epoch 27: Train err: 0.296875, Train loss: 0.5669098161160946 |  
Validation err: 0.317, Validation loss: 0.5996565371751785  
Epoch 28: Train err: 0.3015, Train loss: 0.5659177647903562 |  
Validation err: 0.318, Validation loss: 0.5953392386436462  
Epoch 29: Train err: 0.296125, Train loss: 0.56122862175107 |  
Validation err: 0.3175, Validation loss: 0.6011282876133919  
Epoch 30: Train err: 0.299375, Train loss: 0.5681498721241951 |  
Validation err: 0.318, Validation loss: 0.5956272780895233  
Finished Training  
Total time elapsed: 82.97 seconds
```





### Part (c) - 2pt

Based on your result from Part(a), suggest another set of hyperparameter values to try. Justify your choice.

```
# Hyperparameters were chosen as:
# network = large_net
# batch_size = 256
# learning_rate = 0.01
# epoch = 25
# The results in Part(a) were good, but there was still slight
# overfitting near
# the end, starting from approximately epoch 21-22.
```

### Part (d) - 1pt

Train the model with the hyperparameters you chose in part(c), and include the training curve.

```
large_net = LargeNet()
train_net(large_net, 256, 0.01, 25)
model_path_large = get_model_name("large", batch_size=256,
```

```
learning_rate=0.01, epoch=24)
plot_training_curve(model_path_large)
```

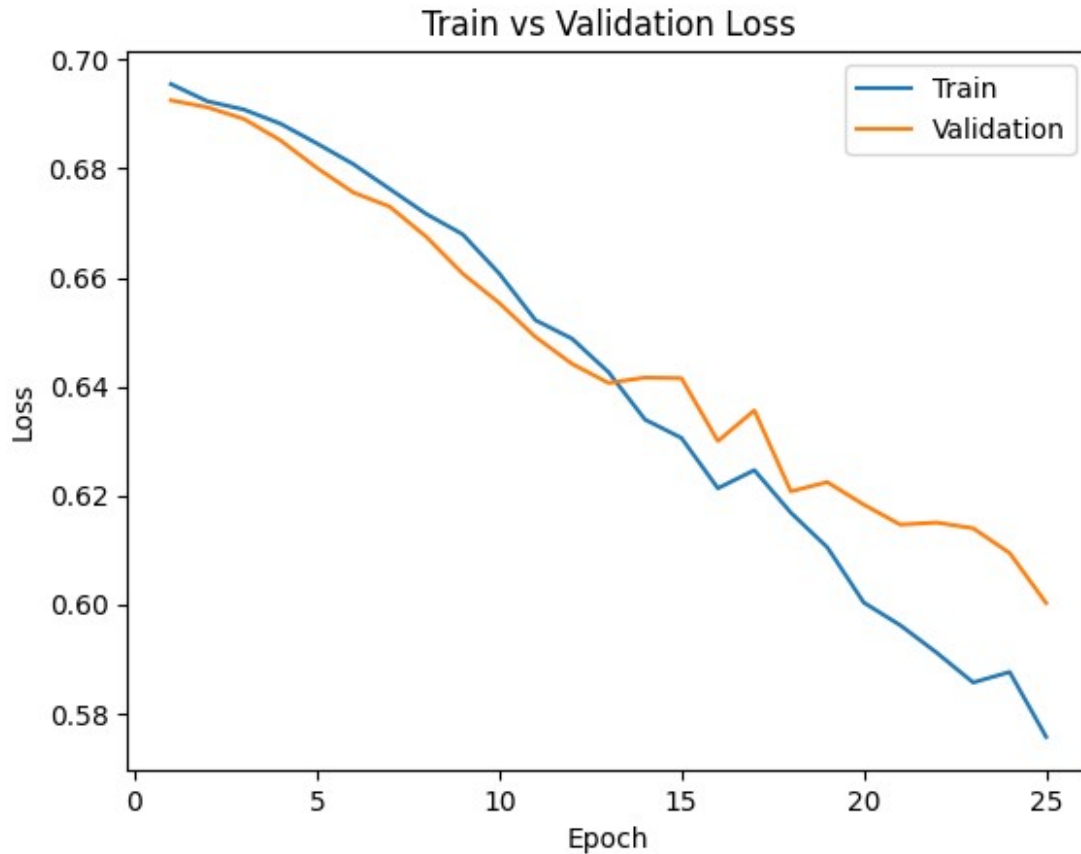
Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.50225, Train loss: 0.6955268159508705 |
Validation err: 0.491, Validation loss: 0.6925421804189682
Epoch 2: Train err: 0.49775, Train loss: 0.6923610251396894 |
Validation err: 0.4655, Validation loss: 0.6912638545036316
Epoch 3: Train err: 0.44675, Train loss: 0.6908448059111834 |
Validation err: 0.4205, Validation loss: 0.6891277506947517
Epoch 4: Train err: 0.437, Train loss: 0.6882675737142563 |Validation
err: 0.421, Validation loss: 0.6852065697312355
Epoch 5: Train err: 0.43725, Train loss: 0.6846398003399372 |
Validation err: 0.4145, Validation loss: 0.6801528334617615
Epoch 6: Train err: 0.431, Train loss: 0.6808352507650852 |Validation
err: 0.4145, Validation loss: 0.6756336018443108
Epoch 7: Train err: 0.421625, Train loss: 0.6762709021568298 |
Validation err: 0.407, Validation loss: 0.6730251237750053
Epoch 8: Train err: 0.405625, Train loss: 0.6716891080141068 |
Validation err: 0.4025, Validation loss: 0.6675112992525101
Epoch 9: Train err: 0.4005, Train loss: 0.6679535750299692 |Validation
err: 0.39, Validation loss: 0.6607388854026794
Epoch 10: Train err: 0.3835, Train loss: 0.6607740316540003 |
Validation err: 0.3985, Validation loss: 0.6553630828857422
Epoch 11: Train err: 0.37425, Train loss: 0.6521737668663263 |
Validation err: 0.376, Validation loss: 0.6490888521075249
Epoch 12: Train err: 0.3795, Train loss: 0.6487875431776047 |
Validation err: 0.3715, Validation loss: 0.6441735848784447
Epoch 13: Train err: 0.357, Train loss: 0.6426474433392286 |Validation
err: 0.3605, Validation loss: 0.6406564712524414
Epoch 14: Train err: 0.358, Train loss: 0.6339474860578775 |Validation
err: 0.373, Validation loss: 0.641641654074192
Epoch 15: Train err: 0.353875, Train loss: 0.6305643077939749 |
Validation err: 0.3705, Validation loss: 0.64152792096138
Epoch 16: Train err: 0.34125, Train loss: 0.6213253419846296 |
Validation err: 0.3565, Validation loss: 0.62999027967453
Epoch 17: Train err: 0.341875, Train loss: 0.6246551889926195 |
Validation err: 0.3615, Validation loss: 0.6356113106012344
Epoch 18: Train err: 0.338625, Train loss: 0.6168491505086422 |
Validation err: 0.3395, Validation loss: 0.6207610294222832
Epoch 19: Train err: 0.333875, Train loss: 0.6104864776134491 |
Validation err: 0.343, Validation loss: 0.6224529445171356
Epoch 20: Train err: 0.324875, Train loss: 0.6003531944006681 |
Validation err: 0.338, Validation loss: 0.6183240413665771
Epoch 21: Train err: 0.321375, Train loss: 0.5961896348744631 |
Validation err: 0.3325, Validation loss: 0.6146606802940369
Epoch 22: Train err: 0.315625, Train loss: 0.5911471806466579 |
Validation err: 0.336, Validation loss: 0.6150308176875114
Epoch 23: Train err: 0.308625, Train loss: 0.5856778360903263 |
```

Validation err: 0.3255, Validation loss: 0.6139866262674332  
Epoch 24: Train err: 0.311625, Train loss: 0.5876005031168461 |  
Validation err: 0.324, Validation loss: 0.6094432324171066  
Epoch 25: Train err: 0.3025, Train loss: 0.5757022220641375 |  
Validation err: 0.316, Validation loss: 0.6002651676535606  
Finished Training  
Total time elapsed: 71.39 seconds





## Part 4. Evaluating the Best Model [15 pt]

### Part (a) - 1pt

Choose the **best** model that you have so far. This means choosing the best model checkpoint, including the choice of `small_net` vs `large_net`, the `batch_size`, `learning_rate`, and the **epoch number**.

Modify the code below to load your chosen set of weights to the model object `net`.

```
net = LargeNet()
train_net(net, 256, 0.01, 25)
model_path = get_model_name(net.name, batch_size=256,
learning_rate=0.01, epoch=24)
state = torch.load(model_path)
net.load_state_dict(state)
```

Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.497125, Train loss: 0.6934096086770296 |  
Validation err: 0.468, Validation loss: 0.6913570985198021

Epoch 2: Train err: 0.453125, Train loss: 0.6905003543943167 |  
Validation err: 0.44, Validation loss: 0.688907079398632



Epoch 3: Train err: 0.445875, Train loss: 0.6875452101230621 |  
Validation err: 0.434, Validation loss: 0.6848585233092308  
Epoch 4: Train err: 0.43575, Train loss: 0.6822220403701067 |  
Validation err: 0.4335, Validation loss: 0.6786830052733421  
Epoch 5: Train err: 0.431375, Train loss: 0.676894897595048 |  
Validation err: 0.417, Validation loss: 0.673255443572998  
Epoch 6: Train err: 0.421375, Train loss: 0.6736960113048553 |  
Validation err: 0.405, Validation loss: 0.66932263225317  
Epoch 7: Train err: 0.405, Train loss: 0.6672013718634844 |Validation  
err: 0.388, Validation loss: 0.6631332784891129  
Epoch 8: Train err: 0.393875, Train loss: 0.6621953304857016 |  
Validation err: 0.3795, Validation loss: 0.6555878296494484  
Epoch 9: Train err: 0.382375, Train loss: 0.6557383760809898 |  
Validation err: 0.37, Validation loss: 0.6476695016026497  
Epoch 10: Train err: 0.368875, Train loss: 0.6490027382969856 |  
Validation err: 0.367, Validation loss: 0.6428123041987419  
Epoch 11: Train err: 0.36475, Train loss: 0.642054982483387 |  
Validation err: 0.368, Validation loss: 0.6401692926883698  
Epoch 12: Train err: 0.36, Train loss: 0.6374363955110312 |Validation  
err: 0.3725, Validation loss: 0.6437254026532173  
Epoch 13: Train err: 0.3525, Train loss: 0.635645866394043 |Validation  
err: 0.3645, Validation loss: 0.6336137726902962  
Epoch 14: Train err: 0.3475, Train loss: 0.6250239498913288 |  
Validation err: 0.3625, Validation loss: 0.6324087083339691  
Epoch 15: Train err: 0.34875, Train loss: 0.6219928823411465 |  
Validation err: 0.3575, Validation loss: 0.6280075907707214  
Epoch 16: Train err: 0.3345, Train loss: 0.6132696438580751 |  
Validation err: 0.351, Validation loss: 0.6226278096437454  
Epoch 17: Train err: 0.332625, Train loss: 0.6122955922037363 |  
Validation err: 0.357, Validation loss: 0.6286576688289642  
Epoch 18: Train err: 0.327875, Train loss: 0.6045868936926126 |  
Validation err: 0.336, Validation loss: 0.6162997633218765  
Epoch 19: Train err: 0.3265, Train loss: 0.6010097619146109 |  
Validation err: 0.3375, Validation loss: 0.6118751764297485  
Epoch 20: Train err: 0.318375, Train loss: 0.591007798910141 |  
Validation err: 0.331, Validation loss: 0.6103062033653259  
Epoch 21: Train err: 0.31275, Train loss: 0.5846749544143677 |  
Validation err: 0.3355, Validation loss: 0.6062348783016205  
Epoch 22: Train err: 0.308125, Train loss: 0.5792908165603876 |  
Validation err: 0.3305, Validation loss: 0.6022427529096603  
Epoch 23: Train err: 0.30575, Train loss: 0.5745287388563156 |  
Validation err: 0.3195, Validation loss: 0.6053818762302399  
Epoch 24: Train err: 0.3035, Train loss: 0.5721510984003544 |  
Validation err: 0.326, Validation loss: 0.5991306826472282  
Epoch 25: Train err: 0.294125, Train loss: 0.5622412748634815 |  
Validation err: 0.3165, Validation loss: 0.5889866054058075  
Finished Training  
Total time elapsed: 71.98 seconds

<All keys matched successfully>

## Part (b) - 2pt

Justify your choice of model from part (a).

*# It resulted in a slightly smaller validation error and loss.*

## Part (c) - 2pt

Using the code in Part 0, any code from lecture notes, or any code that you write, compute and report the **test classification error** for your chosen model.

```
# If you use the `evaluate` function provided in part 0, you will need to
# set batch_size > 1
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=256)
```

```
criterion = nn.BCEWithLogitsLoss()
test_error, test_loss = evaluate(net, test_loader, criterion)
print("Test classification error:", test_error, test_loss)
```

```
Files already downloaded and verified
Files already downloaded and verified
Test classification error: 0.3105 0.5830739140510559
```

## Part (d) - 3pt

How does the test classification error compare with the **validation error**? Explain why you would expect the test error to be *higher* than the validation error.

```
# The test classification error is marginally lower than the
validation error
# The test error is often higher than the validation error because the
model is
# tuned to perform well on the validation set during development, but
the test
# set is completely unseen, providing a more unbiased measure of the
model's
# performance on new data.
```

## Part (e) - 2pt

Why did we only use the test data set at the very end? Why is it important that we use the test data as little as possible?

```
# We use the test dataset at the very end to ensure that the  
evaluation of the  
# model's performance is as unbiased and accurate as possible. It's  
important to  
# use the test data as little as possible to avoid the risk of  
inadvertently  
# tuning the model to this dataset as well, which would compromise its  
ability  
# to generalize to new data.
```

## Part (f) - 5pt

How does your best CNN model compare with a 2-layer ANN model (no convolutional layers) on classifying cat and dog images. You can use a 2-layer ANN architecture similar to what you used in Lab 1. You should explore different hyperparameter settings to determine how well you can do on the validation dataset. Once satisfied with the performance, you may test it out on the test data.

Hint: The ANN in lab 1 was applied on greyscale images. The cat and dog images are colour (RGB) and so you will need to flatten and concatenate all three colour layers before feeding them into an ANN.

```
# 2 layer ANN model:  
import torch  
import torch.nn as nn  
import torch.nn.functional as F  
from torchvision import datasets, transforms  
import matplotlib.pyplot as plt # for plotting  
import torch.optim as optim  
  
torch.manual_seed(1) # set the random seed  
  
class Lab1ANN(nn.Module):  
    def __init__(self):  
        super(Lab1ANN, self).__init__()  
        self.name = "Lab1ANN"  
        self.layer1 = nn.Linear(32 * 32 * 3, 30)  
        self.layer2 = nn.Linear(30, 1)  
    def forward(self, img):  
        flattened = img.view(-1, 32 * 32 * 3)  
        activation1 = self.layer1(flattened)  
        activation1 = F.relu(activation1)  
        activation2 = self.layer2(activation1)  
        activation2 = activation2.squeeze(1)  
        return activation2  
  
Lab1ANN = Lab1ANN()  
train_net(Lab1ANN, batch_size = 256, learning_rate = 0.01, num_epochs  
= 29 )
```

```
criterion = nn.BCEWithLogitsLoss()
test_error, test_loss = evaluate(net, test_loader, criterion)
print("Test classification error and loss:", test_error, test_loss)
```

*# CNN performed better and had smaller errors compared to the ANN model.*

*# Validation err: 0.392, Validation loss: 0.8267480581998825*

Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.421875, Train loss: 0.6742314212024212 |
Validation err: 0.3945, Validation loss: 0.6581188961863518
Epoch 2: Train err: 0.38075, Train loss: 0.650376234203577 |Validation
err: 0.39, Validation loss: 0.6518372967839241
Epoch 3: Train err: 0.371, Train loss: 0.6385549809783697 |Validation
err: 0.38, Validation loss: 0.6489328816533089
Epoch 4: Train err: 0.358875, Train loss: 0.6303329616785049 |
Validation err: 0.3875, Validation loss: 0.6505187824368477
Epoch 5: Train err: 0.353, Train loss: 0.6234526988118887 |Validation
err: 0.379, Validation loss: 0.6453062519431114
Epoch 6: Train err: 0.340125, Train loss: 0.6170223131775856 |
Validation err: 0.379, Validation loss: 0.6477555632591248
Epoch 7: Train err: 0.336625, Train loss: 0.6100924164056778 |
Validation err: 0.374, Validation loss: 0.643535666167736
Epoch 8: Train err: 0.32775, Train loss: 0.6041284091770649 |
Validation err: 0.374, Validation loss: 0.645165205001831
Epoch 9: Train err: 0.32125, Train loss: 0.5959411058574915 |
Validation err: 0.384, Validation loss: 0.6470723003149033
Epoch 10: Train err: 0.312125, Train loss: 0.5877368673682213 |
Validation err: 0.3695, Validation loss: 0.6487203687429428
Epoch 11: Train err: 0.30275, Train loss: 0.5800998397171497 |
Validation err: 0.367, Validation loss: 0.6472310572862625
Epoch 12: Train err: 0.293625, Train loss: 0.5699734576046467 |
Validation err: 0.365, Validation loss: 0.6479431167244911
Epoch 13: Train err: 0.293875, Train loss: 0.5669378507882357 |
Validation err: 0.368, Validation loss: 0.6554356440901756
Epoch 14: Train err: 0.280875, Train loss: 0.5559590179473162 |
Validation err: 0.373, Validation loss: 0.6492954865098
Epoch 15: Train err: 0.269, Train loss: 0.5451897978782654 |Validation
err: 0.369, Validation loss: 0.6529053822159767
Epoch 16: Train err: 0.271625, Train loss: 0.5401697484776378 |
Validation err: 0.374, Validation loss: 0.6571093574166298
Epoch 17: Train err: 0.264375, Train loss: 0.5304668704047799 |
Validation err: 0.3585, Validation loss: 0.6480053588747978
Epoch 18: Train err: 0.248875, Train loss: 0.5231432868167758 |
Validation err: 0.3665, Validation loss: 0.6499830931425095
Epoch 19: Train err: 0.24675, Train loss: 0.509241764433682 |
Validation err: 0.3775, Validation loss: 0.6848082914948463
Epoch 20: Train err: 0.243125, Train loss: 0.5054017417132854 |
```

Validation err: 0.3715, Validation loss: 0.659767858684063  
Epoch 21: Train err: 0.23375, Train loss: 0.49454705603420734 |  
Validation err: 0.3685, Validation loss: 0.6676343679428101  
Epoch 22: Train err: 0.223375, Train loss: 0.4882209366187453 |  
Validation err: 0.371, Validation loss: 0.6717003881931305  
Epoch 23: Train err: 0.2185, Train loss: 0.47743367310613394 |  
Validation err: 0.3685, Validation loss: 0.6796082779765129  
Epoch 24: Train err: 0.210625, Train loss: 0.4681494077667594 |  
Validation err: 0.3695, Validation loss: 0.7015681117773056  
Epoch 25: Train err: 0.215875, Train loss: 0.4645947962999344 |  
Validation err: 0.3575, Validation loss: 0.6696438789367676  
Epoch 26: Train err: 0.20125, Train loss: 0.4485793886706233 |  
Validation err: 0.367, Validation loss: 0.7097871229052544  
Epoch 27: Train err: 0.196625, Train loss: 0.44094559643417597 |  
Validation err: 0.354, Validation loss: 0.6874147802591324  
Epoch 28: Train err: 0.18625, Train loss: 0.4256433192640543 |  
Validation err: 0.3595, Validation loss: 0.7027585431933403  
Epoch 29: Train err: 0.19, Train loss: 0.4227746408432722 |Validation  
err: 0.392, Validation loss: 0.8267480581998825  
Finished Training  
Total time elapsed: 68.07 seconds  
Test classification error and loss: 0.3105 0.5833320766687393